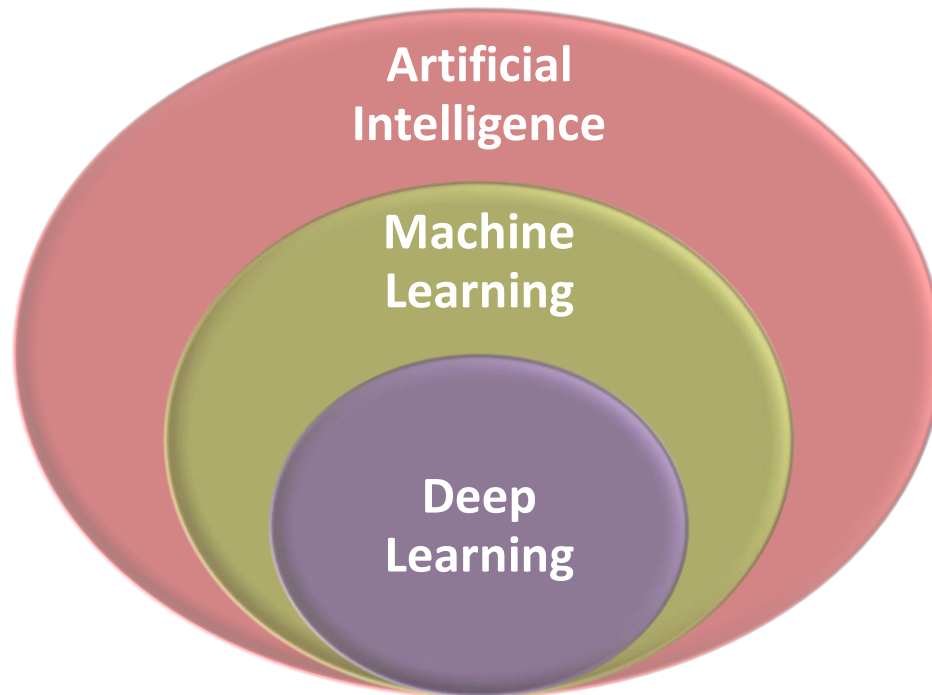


Deep Learning

Deep Learning

- Deep learning is an immensely rich subfield of machine learning, with powerful applications ranging from machine perception to natural language processing, all the way up to creative AI.



Artificial Intelligence, Machine Learning and Deep Learning

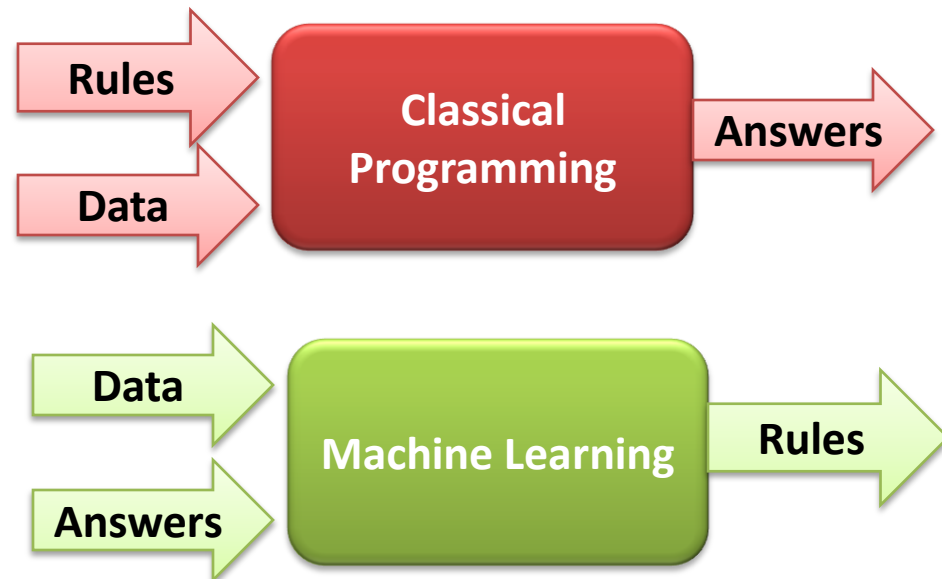
Artificial Intelligence

- A concise definition of **AI** would be:
“the effort to automate intellectual tasks normally performed by humans.”
- AI is a very general field which encompasses machine learning and deep learning, but also includes many more approaches that do not involve any learning.

Machine Learning

- With Machine Learning, humans would input data as well as the answers expected from the data, and out would come the rules.
- These rules could then be applied to new data to produce original answers.

A machine learning system is "trained" rather than explicitly programmed. It is presented with many "examples" relevant to a task, and it finds statistical structure in these examples which eventually allows the system to come up with rules for automating the task.



To do machine learning, we need three things:

- Input data points.
- Examples of the expected output.
- A way to measure if the algorithm is doing a good job, to measure the distance between its current output and its expected output.

This is used as a feedback signal to adjust the way the algorithm works.

This adjustment step is what we call "learning".

- **Machine learning is, technically:**

Searching for useful representations of some input data, within a pre-defined space of possibilities, using guidance from some feedback signal. This simple idea allows for solving a remarkably broad range of intellectual tasks, from speech recognition to autonomous car driving.

Deep Learning


- Deep learning is a specific subfield of machine learning, a new take on learning representations from data which puts an emphasis on learning successive "layers" of increasingly meaningful representations.
- The "deep" in "deep learning" simply stands for this idea of successive layers of representations—how many layers contribute to a model of the data is called the "depth" of the model.

Deep Learning



- Modern deep learning often involves tens or even hundreds of successive layers of representation and they are all learned automatically from exposure to training data.
- Other approaches to machine learning tend to focus on learning only one or two layers of representation of the data, sometimes called **"shallow learning"**.

Deep Learning

- In deep learning, these layered representations are learned via models called "neural networks", structured in literal layers stacked one after the other.
- The term "neural network" is a reference to neurobiology.




Deep Neural Networks




Understanding how deep learning works

- Deep neural networks do input-to-target mapping via a deep sequence of simple data transformations called "layers", and that these data transformations are learned by exposure to examples.

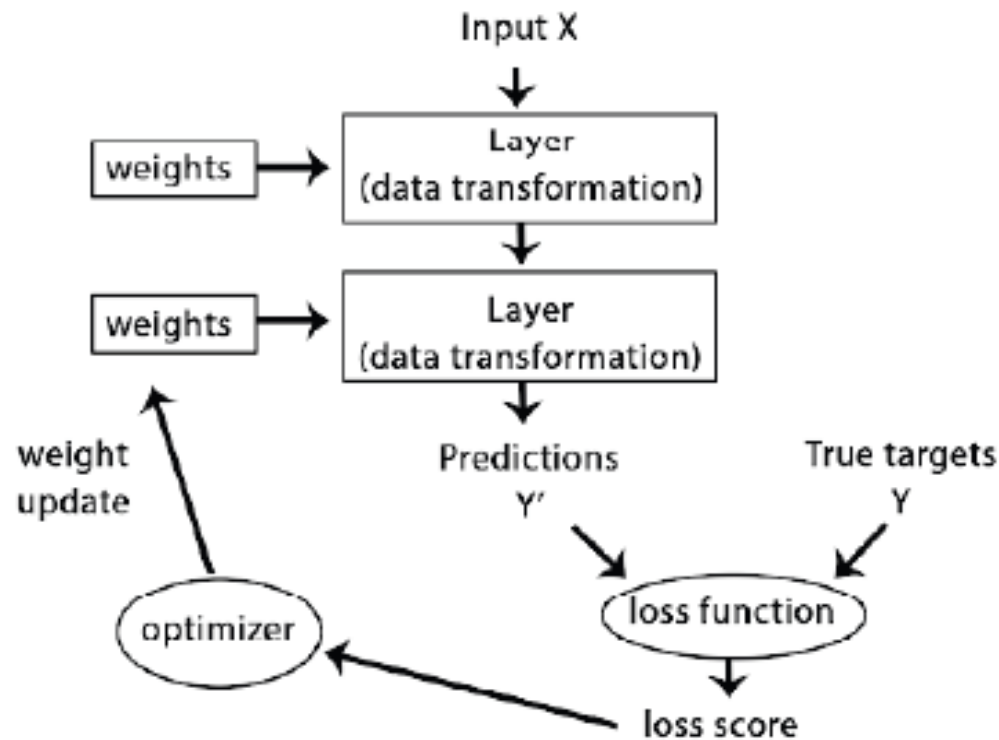
- The specification of what a layer does to its input data is stored in the layer's "**weights**", which in essence are a bunch of numbers.
- Weights are also sometimes called the "**parameters**" of a layer.
- "**learning**" will mean finding a set of values for the weights of all layers in a network, such that the network will correctly map your example inputs to their associated targets.



- To control the output of a neural network, you need to be able to measure how far this output is from what you expected.
- This is the job of the "**loss function**" of the network, also called "**objective function**".
- The loss function takes the predictions of the network and the true target , and computes a distance score, capturing how well the network has done on this specific example.



- The fundamental trick in deep learning is to use the loss score as a feedback signal to adjust the value of the weights by a little bit, in a direction that would lower the loss score for the current example.
- This adjustment is the job of the "optimizer", which implements what is called the "**backpropagation**" algorithm, the central algorithm in deep learning.



What deep learning has achieved so far

- Deep learning has achieved the following breakthroughs, all in historically difficult areas of machine learning:
 - ✓ *Near-human level image classification.*
 - ✓ *Near-human level speech recognition.*
 - ✓ *Near-human level handwriting transcription.*
 - ✓ *Improved text-to-speech conversion.*
 - ✓ *Digital assistants such as Amazon Alexa.*
 - ✓ *Improved ad targeting, as used by Google, Bing.*
 - ✓ *Improved search results on the web.*


What makes deep learning different

- Deep learning took off so quickly is primarily that it offered better performance on many problems.
- Deep learning is also making problem-solving much easier, because it completely automates what used to be the most crucial step in a machine learning workflow:
"feature engineering".

Human Brain

Cerebrum

**Recurrent
Neural Network**



**Convolution
Neural Network**

**Artificial
Neural Network**

Input Layer

- This layer accepts input features.
- It provides information from the outside world to the network, no computation is performed at this layer, nodes here just pass on the information(features) to the hidden layer.


Hidden Layer

- Nodes of this layer are not exposed to the outer world, they are the part of the abstraction provided by any neural network.
- Hidden layer performs all sort of computation on the features entered through the input layer and transfer the result to the output layer.

Output Layer

- This layer returns the final output computed by the network to the application.


Neural Network with Single Neuron




$$\hat{y} = x_0 w_0 + x_1 w_1 + x_2 w_2 = \sum_i x_i w_i$$

$$\begin{bmatrix} w_0 & w_1 & w_2 \end{bmatrix} \begin{bmatrix} x_0 \\ x_1 \\ x_2 \end{bmatrix}$$


Multi-layer Perceptron




Linear Function



$$y_0 = x_0 w_0 w_1 w_2 = x_0 w_C$$





Non-Linear Function



$$y_0 = \sigma(\sigma(\sigma(x_0 w_0) w_1) w_2)$$

Simple Example



$$\mathbf{y_pred} = \mathbf{xw}$$

$$\mathbf{E} = (\hat{\mathbf{y}} - \mathbf{y})^2 = (\mathbf{xw} - \mathbf{y})^2 \quad \textit{sum of squared errors}$$

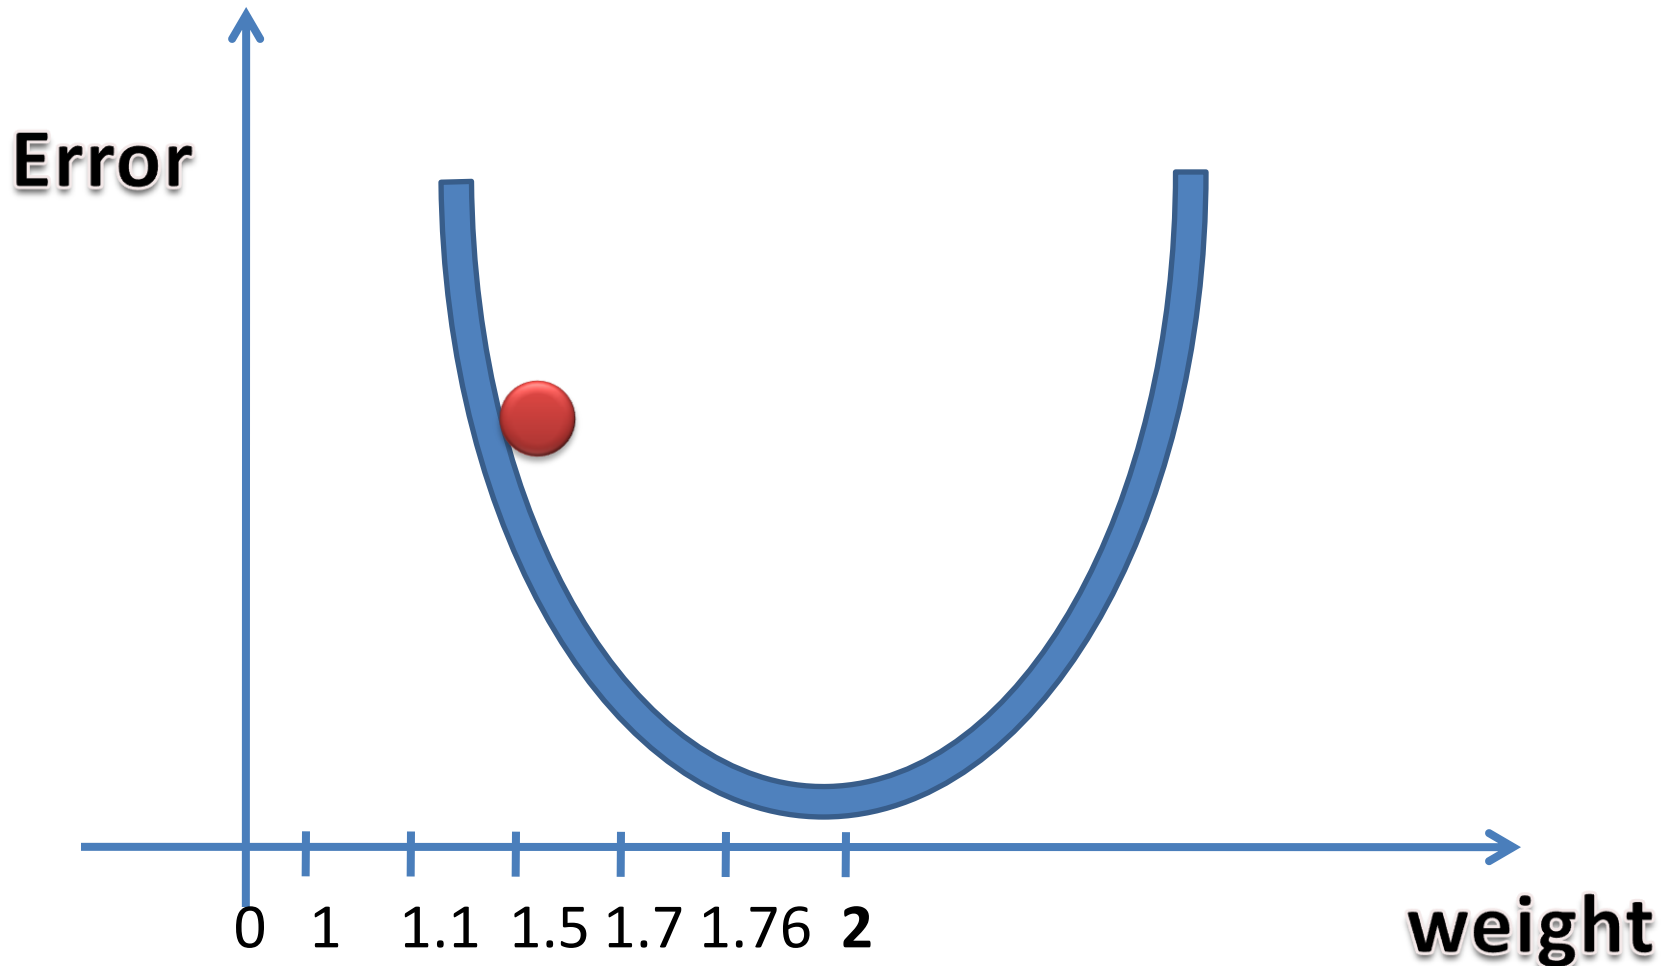
$$\frac{\partial \mathbf{E}}{\partial \mathbf{w}} = 2\mathbf{x}(\mathbf{xw} - \mathbf{y}) \quad \text{Derivative}$$

$$\mathbf{w} - \alpha \frac{\partial \mathbf{E}}{\partial \mathbf{w}} = \mathbf{w} - \alpha 2\mathbf{x}(\mathbf{xw} - \mathbf{y}) \quad \text{Update Rule}$$

- Random weight initialization $w=0.5$
- learning rating $\alpha = 0.1$
- $w_{\text{new}} = w - 0.1 * 2x(xw-y)$
- (x,y)
- $(2,4) \ w=0.5 \leftarrow 0.5 - 0.1*2*2(2*0.5-4) = 1.7$
- $(1,2) \ w = 1.7 \leftarrow 1.7-0.1*2*1(1*1.7-2) = 1.76$
- $(3,6) \ w= 1.76 \leftarrow 1.76-0.1*2*3(3*1.76-6)=2.192$
- $w \sim 2$

Oversimplified Gradient Descent:

- Calculate slope at current position
- If slope is negative, move right
- If slope is positive, move left
- (Repeat until slope == 0)



Convolution Neural Network

- A convolutional neural network (CNN, or ConvNet) is a class of deep neural networks, most commonly applied to analyzing visual imagery.

Convolution Neural Network

- Facebook uses neural nets for their automatic tagging algorithms
- Google for their photo search
- Amazon for their product recommendations
- Pinterest for their home feed personalization
- Instagram for their search infrastructure.



Image classification

- Most popular, use case of these networks is for image processing.

How Humans identify images??

- For humans, this task of recognition is one of the first skills we learn from the moment we are born and is one that comes naturally and effortlessly as adults.
- Most of the time we are able to immediately characterize the scene and give each object a label, all without even consciously noticing.

**What's
that?**

A Train



**What's
this?**

A Train




Teach Machine


Human skills of being able to :

- Quickly recognize patterns
- Generalize from prior knowledge

illusions



illusions



Convolution Neural Network

Input image



CNN

output image

Happy

Input image



CNN

output image

Sad

How Does Computer See an image??

When a computer sees an image ,it will see an array of pixel values. Depending on the resolution and size of the image, it will see a 32 x 32 x 3 array of numbers (The 3 refers to RGB values).



What we see

08	02	22	97	38	15	00	40	00	75	04	05	07	78	52	12	50	77	91	08
49	49	99	40	17	81	18	57	60	87	17	40	98	43	69	48	04	56	62	00
81	49	31	73	55	79	14	29	93	71	40	67	53	88	30	05	49	13	36	65
52	70	95	23	04	60	11	42	69	24	68	56	01	32	56	71	37	02	36	91
22	31	16	71	51	67	63	89	41	92	36	54	22	40	40	28	66	33	13	80
24	47	32	60	99	03	45	02	44	75	33	53	78	36	84	20	35	17	12	50
32	98	81	28	64	23	67	10	26	38	40	67	59	54	70	66	18	38	64	70
67	26	20	68	02	62	12	20	95	63	94	39	63	08	40	91	66	49	94	21
24	55	58	05	66	73	99	26	97	17	78	78	96	83	14	88	34	89	63	72
21	36	23	09	75	00	76	44	20	45	35	14	00	61	33	97	34	31	33	95
78	17	53	28	22	75	31	67	15	94	03	80	04	62	16	14	09	53	56	92
16	39	05	42	96	35	31	47	55	58	88	24	00	17	54	24	36	29	85	57
86	56	00	48	35	71	89	07	05	44	44	37	44	60	21	58	51	54	17	58
19	80	81	68	05	94	47	69	28	73	92	13	86	52	17	77	04	89	55	40
04	52	08	83	97	35	99	16	07	97	57	32	16	26	26	79	33	27	95	66
88	36	68	87	57	62	20	72	03	46	33	67	46	55	12	32	63	93	53	69
04	42	16	73	38	25	39	11	24	94	72	18	08	46	29	32	40	62	76	36
20	69	36	41	72	30	23	88	34	62	99	69	82	67	59	85	74	04	36	16
20	73	35	29	78	31	90	01	74	31	49	71	48	86	81	16	23	57	05	54
01	70	54	71	83	51	54	69	16	92	33	48	61	43	52	01	89	19	67	48

What computers see

How Does Computer See an image??



What we see

```
08 02 22 97 38 15 00 40 00 75 04 05 07 78 52 12 50 77 91 08
49 49 99 40 17 81 18 37 60 87 17 40 98 43 69 48 04 36 62 00
81 49 31 73 55 79 14 29 93 71 40 67 53 88 30 03 49 13 36 65
52 70 95 23 04 60 11 42 69 24 68 56 01 32 56 71 37 02 36 91
22 31 16 71 51 67 63 89 41 92 36 54 22 40 40 28 66 33 13 80
24 47 32 60 99 03 45 02 44 75 33 53 78 36 84 20 35 17 12 50
32 98 81 28 64 23 67 10 26 38 40 67 59 54 70 66 18 38 64 70
67 26 20 68 02 62 12 20 95 63 94 39 63 08 40 91 66 49 94 21
24 55 58 05 66 73 99 26 97 17 78 78 96 83 14 88 34 89 63 72
21 36 23 09 75 00 76 44 20 45 35 14 00 61 33 97 34 31 33 95
78 17 53 28 22 75 31 67 15 94 03 80 04 62 16 14 09 53 56 92
16 39 05 42 96 35 31 47 55 58 88 24 00 17 54 24 36 29 85 57
86 56 00 48 35 71 89 07 05 44 44 37 44 60 21 58 51 54 17 58
19 80 81 68 05 94 47 69 28 73 92 13 86 52 17 77 04 89 55 40
04 52 08 83 97 35 99 16 07 97 57 32 16 26 26 79 33 27 98 66
88 36 68 87 57 62 20 72 03 46 33 67 46 55 12 32 63 93 53 69
04 42 16 73 38 25 39 11 24 94 72 18 08 46 29 32 40 62 76 36
20 69 36 41 72 30 23 88 34 62 99 69 82 67 59 85 74 04 36 16
20 73 35 29 78 31 90 01 74 31 49 71 48 86 81 16 23 57 05 54
01 70 54 71 83 51 54 69 16 92 33 48 61 43 52 01 89 19 67 48
```

What computers see

These numbers, when we perform image classification, are the only inputs available to the computer.

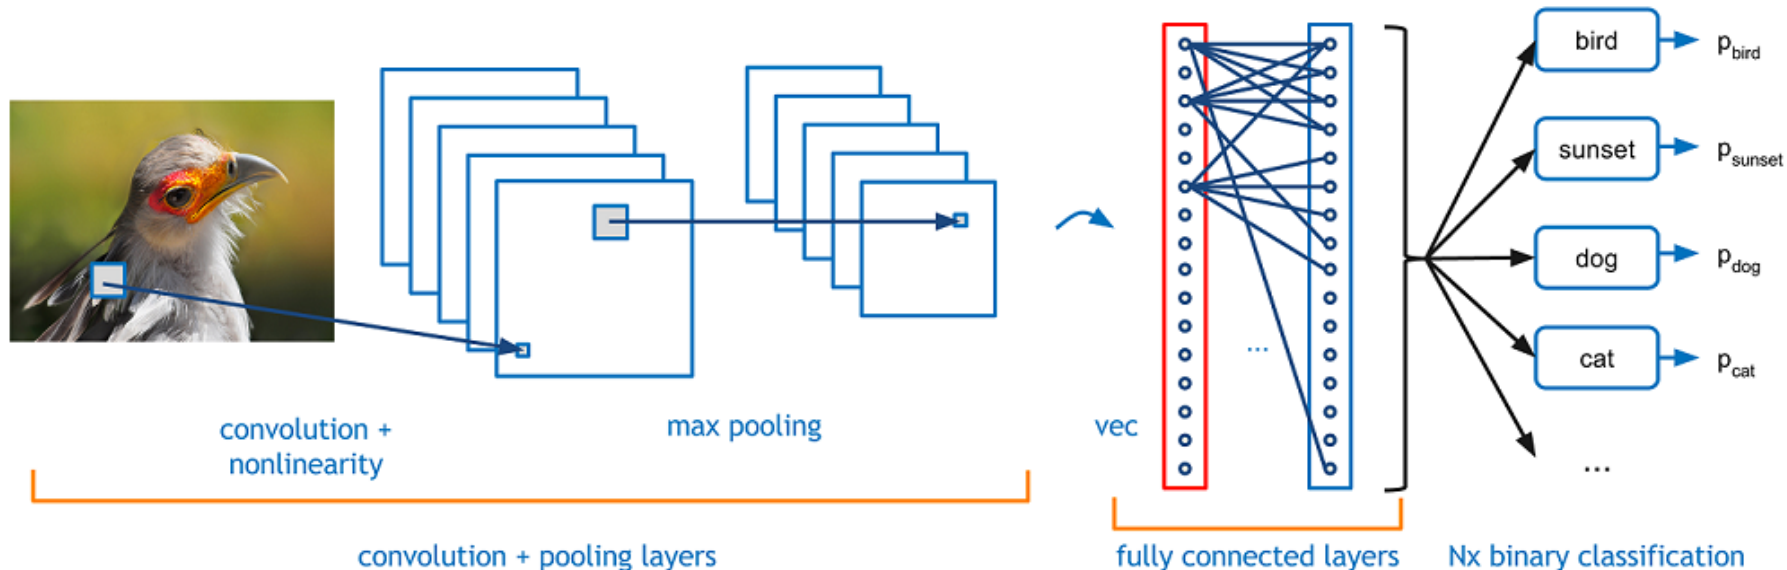
The idea is that you give the computer this array of numbers and it will output numbers that describe the probability of the image being a certain class (80% for flower, 15% for sky etc).

What We Want the Computer to Do

- we want the computer to do is to be able to differentiate between all the images it's given and figure out the **unique features** that make a dog a dog , that make a cat a cat or that make a flower a flower .
- The computer is able to perform image classification by looking for low level features such as edges and curves, and then building up to more abstract concepts through a series of **convolutional layers**.

Convolution Neural Network

CNNs take the image, pass it through a series of convolutional, nonlinear, pooling (downsampling), and fully connected layers, and get an output.



How computer sees an image



Humans see this

-1	-1	-1	-1	-1	-1	-1	-1	-1	0.9	-0	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	0.3	1	0.3	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-0	1	1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	0.8	1	0.6	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	0.5	1	0.8	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	0.1	1	0.9	-0	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-0	1	1	-0	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	0.9	1	0.3	-1	-1	-1	-1	0.5	1	0.9	0.1	-1	-1
-1	-1	0.3	1	0.9	-1	-1	-1	0.1	1	1	1	1	1	-1	-1
-1	-1	0.8	1	0.3	-1	-1	0.4	1	0.7	-0	-0	1	1	-1	-1
-1	-1	1	1	0.1	-1	0.1	1	0.3	-1	-1	-0	1	0.6	-1	-1
-1	-1	1	1	0.8	0.3	1	0.7	-1	-1	-1	0.5	1	0	-1	-1
-1	-1	0.8	1	1	1	1	0.5	0.2	0.8	0.8	1	0.9	-1	-1	-1
-1	-1	-0	0.8	1	1	1	1	1	1	1	1	0.1	-1	-1	-1
-1	-1	-1	-0	0.8	1	1	1	1	1	1	0.2	-1	-1	-1	-1
-1	-1	-1	-1	-1	-0	0.3	0.8	1	0.5	-0	-1	-1	-1	-1	-1

Computer sees this

How computer sees an image

-1	-1	-1	-1	-1	-1	-1	-1	0.9	-0	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	0.3	1	0.3	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-0	1	1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	0.8	1	0.6	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	0.5	1	0.8	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	0.1	1	0.9	-0	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-0	1	1	-0	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	0.9	1	0.3	-1	-1	-1	-1	0.5	1	0.9	0.1	-1	-1
-1	-1	0.3	1	0.9	-1	-1	-1	0.1	1	1	1	1	1	-1	-1
-1	-1	0.8	1	0.3	-1	-1	0.4	1	0.7	-0	-0	1	1	-1	-1
-1	-1	1	1	0.1	-1	0.1	1	0.3	-1	-1	-0	1	0.6	-1	-1
-1	-1	1	1	0.8	0.3	1	0.7	-1	-1	-1	0.5	1	0	-1	-1
-1	-1	0.8	1	1	1	1	0.5	0.2	0.8	0.8	1	0.9	-1	-1	-1
-1	-1	-0	0.8	1	1	1	1	1	1	1	0.1	-1	-1	-1	-1
-1	-1	-1	-0	0.8	1	1	1	1	1	0.2	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-0	0.3	0.8	1	0.5	-0	-1	-1	-1	-1	-1

Computer sees this


-1	-1	-1	-1	-1	-1	-1	-1	-1	0.9	-0	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	0.3		0.3	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-0				-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	0.8		0.6	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	0.5		0.8		-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	0.1					-0	-1	-1	-1	-1	-1	-1
-1	-1	-1	-0					-0	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1							-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	0.9				0.3	-1	-1	-1	-1	0.5		0.1	-1
-1	-1	0.3						-1	-1	0.1					-1
-1	-1	0.8					0.3	-1	-1	0.4		0.7	-0	-0	-1
-1	-1						0.1	-1	0.1		0.3	-1	-1	-0	0.6
-1	-1						0.8	0.3		0.7	-1	-1	-1	0.5	0
-1	-1	0.8							0.5	0.2	0.8	0.8		0.9	-1
-1	-1	-0	0.8											0.1	-1
-1	-1	-1	-0	0.8										0.2	-1
-1	-1	-1	-1	-1	-1	-1	-0	0.3	0.8		0.5	-0	-1	-1	-1

Same matrix, highlight the cells based on cell value

How computer sees an image



Human Vision



Computer Vision

Kernel Matrix examples

1	1	1
1	1	1
1	1	1

Unweighted 3x3 smoothing kernel

0	1	0
1	4	1
0	1	0

Weighted 3x3 smoothing kernel with Gaussian blur

0	-1	0
-1	5	-1
0	-1	0

Kernel to make image sharper

-1	-1	-1
-1	9	-1
-1	-1	-1

Intensified sharper image



Gaussian Blur



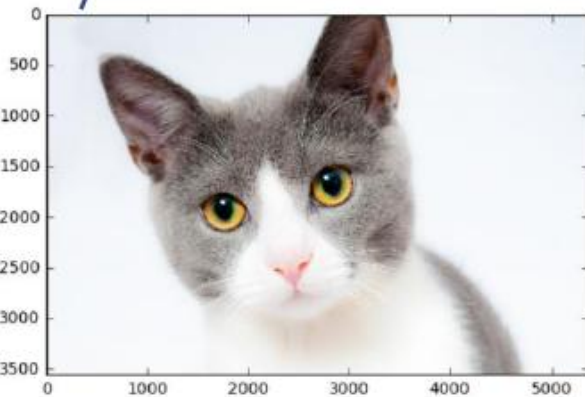


Sharpened image

Convolved features

Different filters capture different features

A filter to capture overall shape and edges

A filter to ignore white spaces on image



The depth in convolution layer

- Every filter gives us a resultant matrix (activation map)

0	0	0	0	0	0	0	0	0
0	1	1	1	0	0	1	1	0
0	1	1	0	1	0	1	1	0
0	1	0	0	1	1	1	0	0
0	0	0	0	0	0	1	1	0
0	1	1	0	0	0	0	1	0
0	1	0	0	1	1	1	0	0
0	1	1	1	0	0	1	1	0
0	0	0	0	0	0	0	0	0


1	1	1
1	1	1
1	1	1

4	5	4	2	3	4	4
5	6	5	4	6	6	5
3	3	3	3	6	6	5
3	3	2	2	4	5	4
3	3	2	2	4	5	4
5	6	4	3	4	5	4
3	4	3	3	4	4	3

The depth in convolution layer

- If we apply 10 different filters then we will get 10 resultant matrices. The depth is 10

0	0	0	0	0	0	0	0	0
0	1	1	1	0	0	1	1	0
0	1	1	0	1	0	1	1	0
0	1	0	0	1	1	1	0	0
0	0	0	0	0	0	1	1	0
0	1	1	0	0	0	0	1	0
0	1	0	0	1	1	1	0	0
0	1	1	1	0	0	1	1	0
0	0	0	0	0	0	0	0	0



3	4	3	2	3	3	3
4	5	5	3	6	5	4
2	3	3	2	5	5	5
3	3	2	2	4	4	3
2	2	2	2	4	5	3
4	6	4	2	3	4	4
2	3	2	3	4	3	2

Max Pooling

Pooling layer

- Pooling helps in further downsizing the data
- Pooling is used to avoid overfitting and increase the robustness
- Reduces lot of computation time

What happens in pooling layer

- Down sampling of the data.
- Since we preserved lot of information in each convolution layer, we can comfortably down sample it, in this pooling layer
- Yes, we do loose some information, but we will not loose the overall integrity of the data
- It is still good enough for a classification model
- Generally we try max pooling with 2X2 filter with stride 2

How many parameters in pooling layer

- No parameters
- We just perform down sampling, that's it. No further activation and no further parameters
- Number of parameters in pooling layer=0

Max Pooling

0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1

Feature Map




1	1	0
4	2	1
0	2	1

**Pool
Feature Map**

Max pooling is usually done with 2x2 windows and stride 2, so as to downsample the feature maps

Max Pooling



Flattening

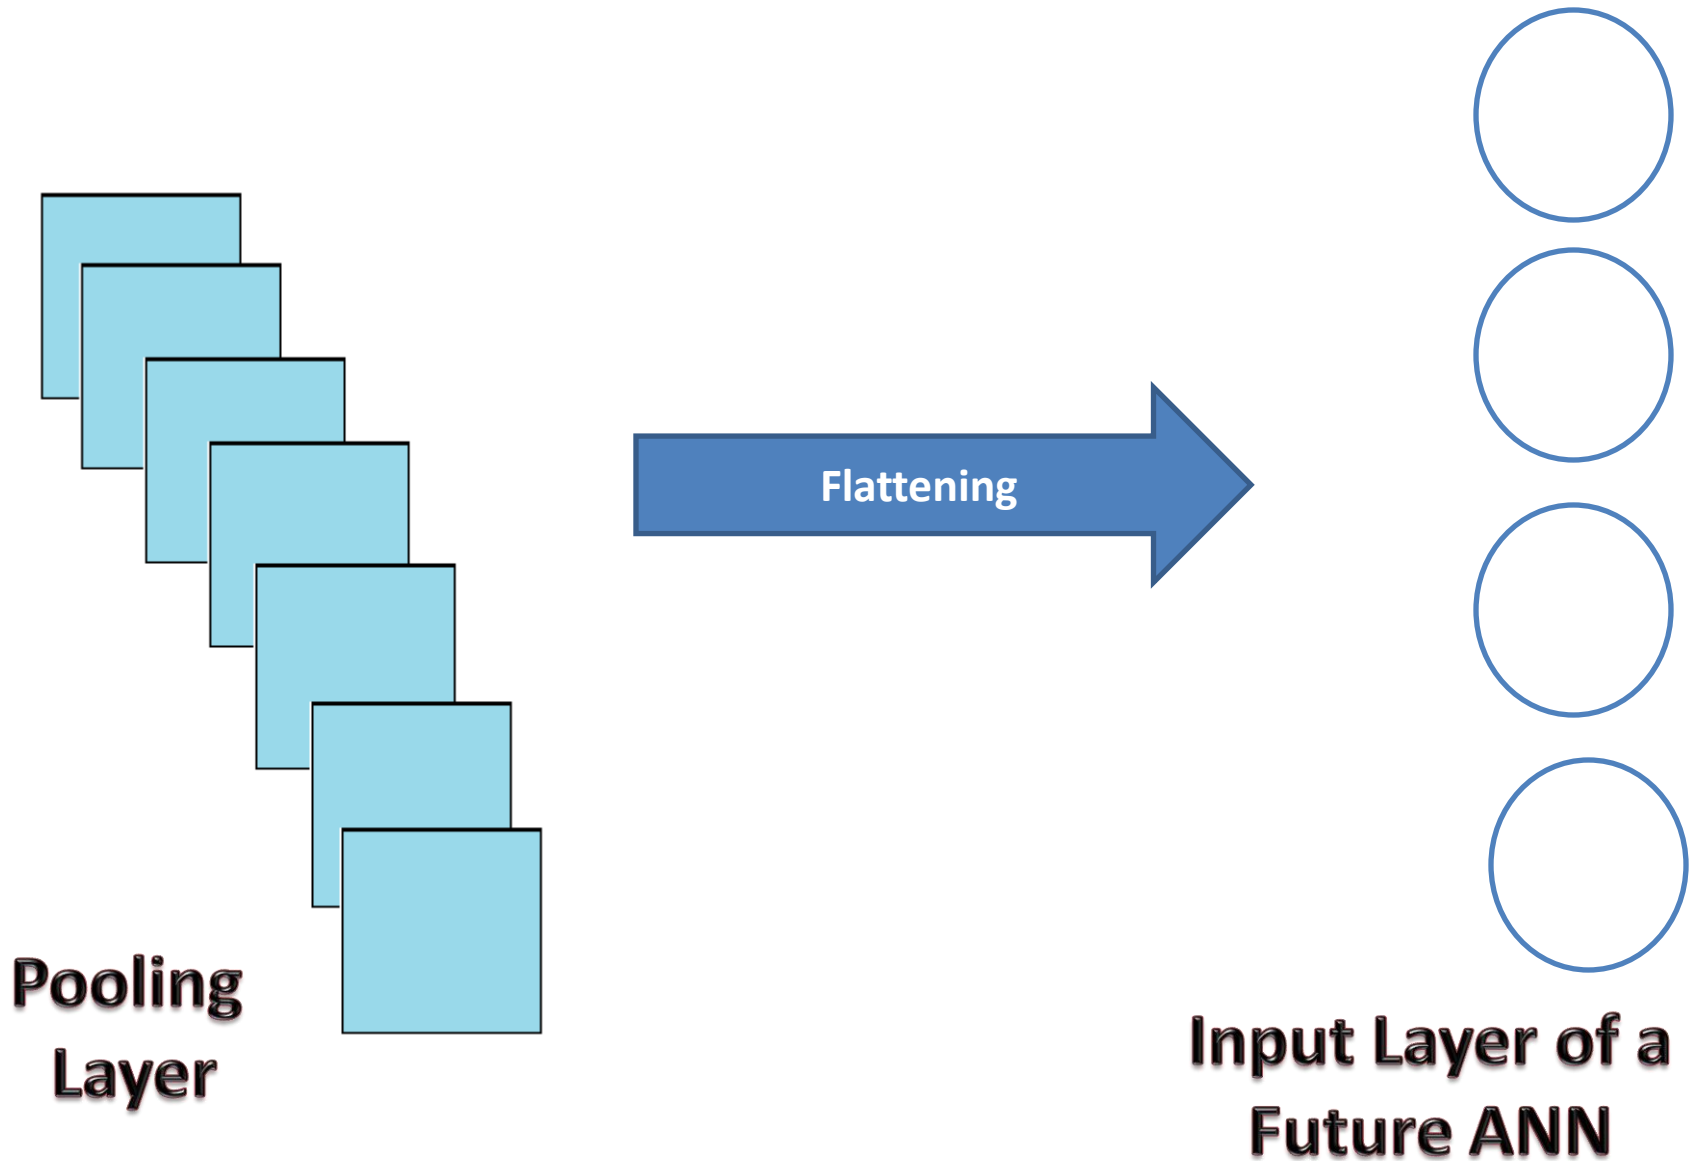
1	1	0
4	2	1
0	2	1

**Pool
Feature Map**




1
1
0
4
2
1
0
2
1

Flattening



Flattening



Full Connection

